

PoC: LLM-driven Ontology

The PoC covers:

- **Structured and unstructured** illustrative **data** about products and logistics.
- **LLM agents** that generate **ontologies** and **retrieve** data to answer users' questions.

*The PoC presents how I would approach some issues of
Ontologies-LLM integration in a (simplified) **edge-case scenario**.*

How to Build Logic Models That Handle Complex Data?

LLM can **generate** consistent **logic models**, also when the *prior knowledge* (e.g., documents and prompts) is **qualitative** and **partial**.

LLM can overcome issues of traditional *expert systems* since hand-made data models often are unreliable and hard to maintain/extend.

The PoC stresses this aspect by providing limited prior knowledge to LLM prompts.

However, when available, more detailed information about the scenario and tasks will likely increase accuracy.

How Can LLM Leverage Ontologies to Retrieve Data?

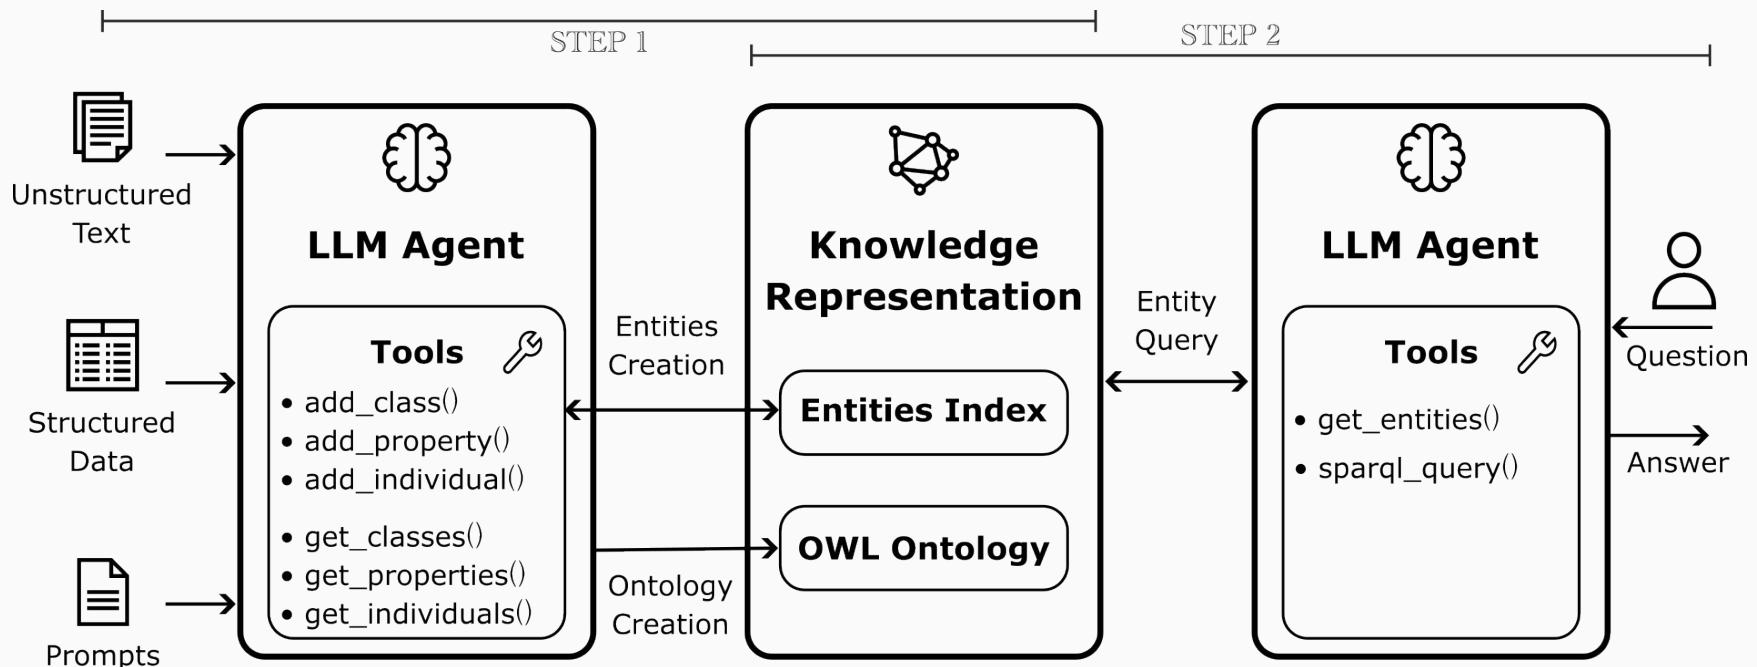
Highly linked data is well-suited for being contextualised in an **ontology** that the **LLM can explore** and benefit from (e.g., via SPARQL queries).

At the cost of *system* complexity, ontologies have several benefits. For example, they support sparse data retrieval, explainability, context-based agent orchestration, consistency checking and error recovery.

The PoC stresses this aspect by enabling the LLM to retrieve data through SPARQL instead of vector search.

However, a combination of vector-search and queries would increase accuracy and decrease computation time.

Implementation



Results

Step 1: Ontology Creation

- The LLM represented **structured** data in the ontology **better** than **unstructured** text.
- The LLM **annotated** ontological entities to enhance **reasoning based on natural language**.
- The generated ontology (about 900 facts) had few **redundancies** and was **human-readable** (i.e., suitable for explainability and maintainability).

Step 2: Data Retrieval

- The LLM explored the ontology and answers were found with **chains of queries**.
- Over **20 questions**: 16 were correct, 1 was partially correct, 1 was incorrect, and 3 had no answer.
- The LLM explained the performed queries in a **reproducible** manner.

Limitations and Improvements

The PoC has a large room for improvement, especially when addressing these limitations:

- The implemented LLM tools address a **fragment of the OWL formalism**, and they do not allow for automatic instance classification.
- **OWL Reasoners** have not been used to support LLM. For example, they could provide **consistency feedbacks** to the LLM for **self-improvement** and **error recovery**.
- LLM with limited prior knowledge generates **granular ontologies**. More **prior knowledge** and a **fine-tuned LLM** would **decrease the number of facts** in the ontology with no semantic loss.
- The PoC does not implement **concurrent LLM agents**, which would decrease computation time.
- Retrieving data only with SPARQL is an extreme case. **Ontologies integrated with Graph or Vector DBs** would improve data retrieval performance.

Thanks for your attention

More details are available in the PoC repository:

https://github.com/buoncubi/PoC_OWL-LLM_Integration

Luca Buoncompagni

06/11/2025